

CISC3024 Pattern Recognition Project

by Zhang Huakang D-B9-2760-6

The notebook file of this project has been published on [kaggle](#).

| Member Name | Contribution Percentage |
|---------------|-------------------------|
| ZHANG HUAKANG | 100% |

1. Data Loading

In this part, the image will be loaded from disk with its corresponding labels. After loading images and labels, a customized class `SatelliteDataset` inherited from `torch.utils.data.Dataset` is used to store this data. When `__getitem__()` function is called, processed images and labels are returned.

```
# Image processing for training data
train_tran = transforms.Compose([
    transforms.RandomRotation(5),
    transforms.RandomHorizontalFlip(0.5),
    transforms.RandomCrop(64, padding=2),
    transforms.ToTensor(),
])

# Image processing for testing and val. data
tran = transforms.Compose([
    transforms.ToTensor(),
])
```

```
# SatelliteDataset
class SatelliteDataset(Dataset):
    def __init__(self, images, labels, is_train):
        self.images_list=images
        self.labels_list=labels
        self.is_train=is_train

    def __len__(self):
        return len(self.labels_list)

    def __getitem__(self, idx):
        # Image Processing
        if self.is_train:
            return train_tran(self.images_list[idx]).to(device),
            self.labels_list[idx]
        return tran(self.images_list[idx]).to(device),
        self.labels_list[idx]
```


Then, `train`, `test` and `val` data will be encapsulated in `torch.utils.data.DataLoader` object to make it easier to get data when training the model.

```
train_iterator = data.DataLoader(SatelliteDataset(X_train,y_train, True),
                                batch_size=BATCH_SIZE)
```


Example Image

Here is some example of dataset.

Unprocessed:

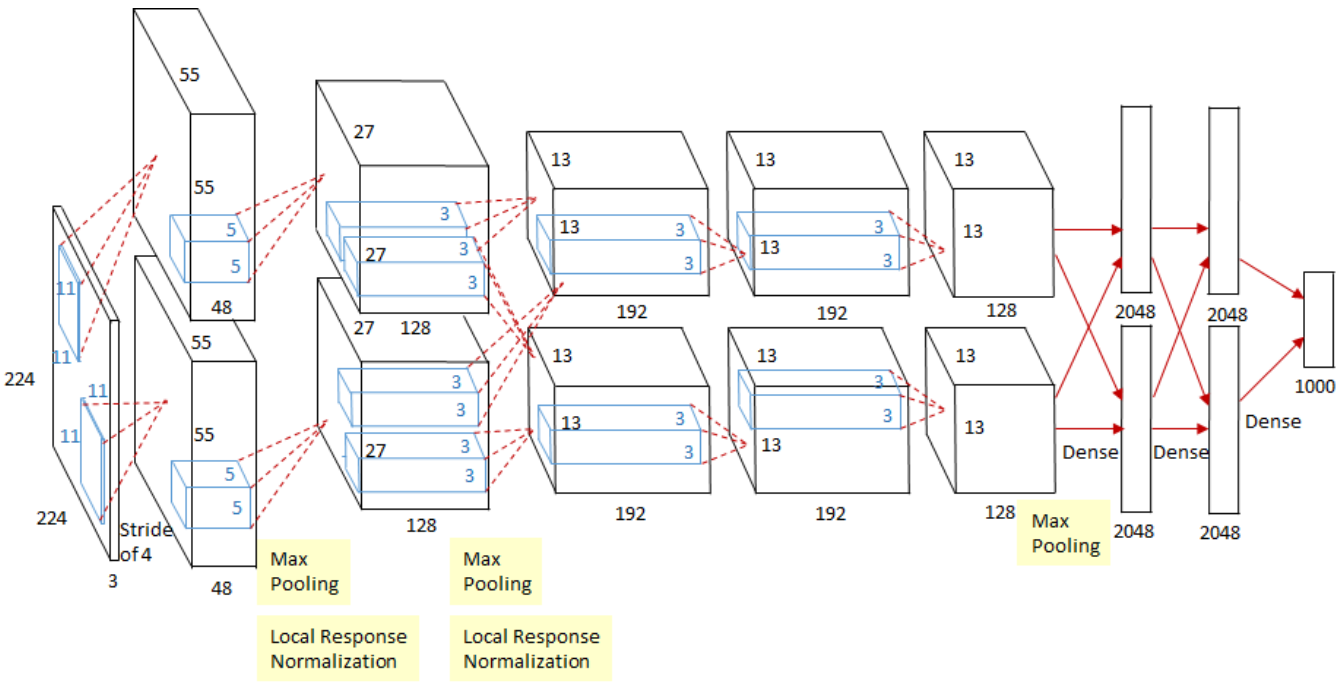


Normalization:



Model Design

AlexNet



code version:

```
class AlexNet(nn.Module):
    def __init__(self, output_dim):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 64, 3, 2, 1),
            nn.MaxPool2d(2),
            nn.ReLU(inplace=True),
            nn.Conv2d(64, 192, 3, padding=1),
            nn.MaxPool2d(2),
            nn.ReLU(inplace=True),
            nn.Conv2d(192, 384, 3, padding=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(384, 256, 3, padding=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(256, 256, 3, padding=1),
            nn.MaxPool2d(2),
            nn.ReLU(inplace=True)
        )
        self.classifier = nn.Sequential(
            nn.Dropout(0.5),
            nn.Linear(4096, 4096),
            nn.ReLU(inplace=True),
            nn.Dropout(0.5),
            nn.Linear(4096, 4096),
            nn.ReLU(inplace=True),
            nn.Linear(4096, output_dim),
        )
    def forward(self, x):
        x = self.features(x)
        h = x.view(x.shape[0], -1)
        x = self.classifier(h)
        return x, h
```

Model Training

Find the best learning rate

In this part, I build a class `LRFinder` to find the best learning rate. In the function `range_test()`,

```
def range_test(self, iterator, end_lr=10, num_iter=100,
               smooth_f=0.05, diverge_th=5):
    lrs = []
    losses = []
    best_loss = float('inf')
    lr_scheduler = ExponentialLR(self.optimizer, end_lr, num_iter)
    iterator = IteratorWrapper(iterator)
    for iteration in range(num_iter):
        loss = self._train_batch(iterator)
```

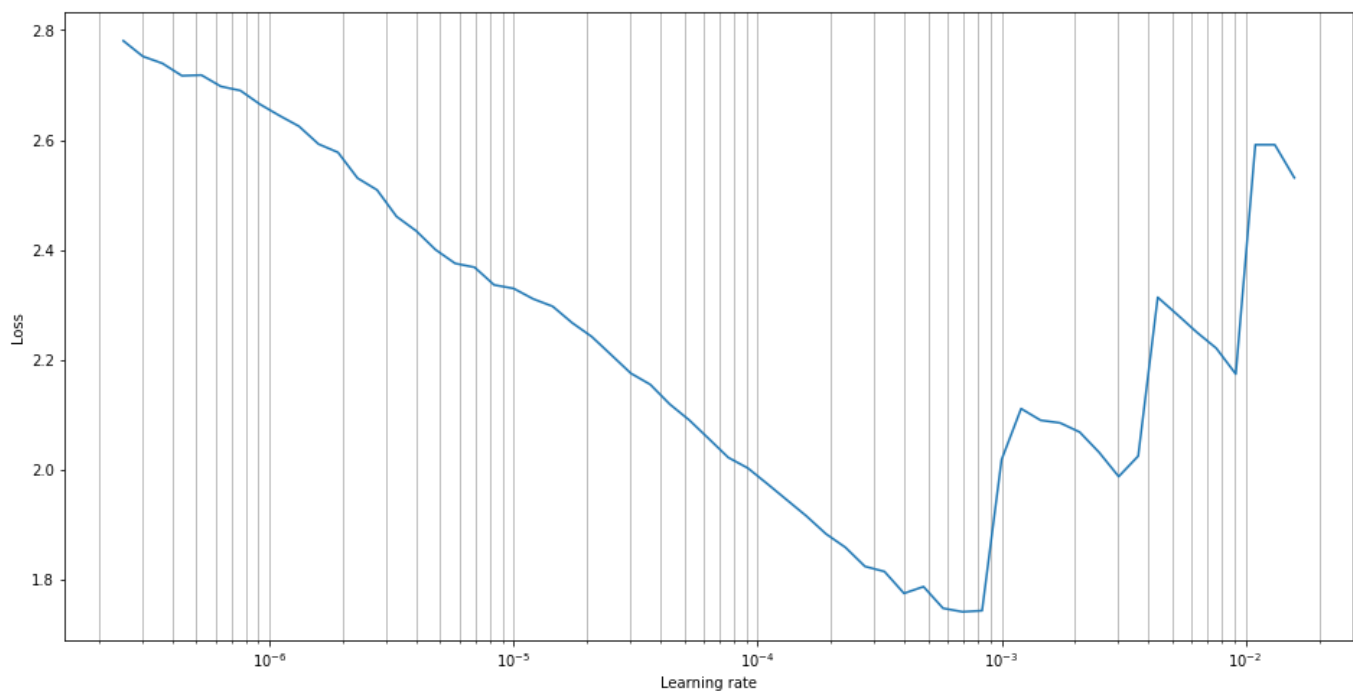
```

    lrs.append(lr_scheduler.get_last_lr()[0])
    # update lr
    lr_scheduler.step()
    if iteration > 0:
        loss = smooth_f * loss + (1 - smooth_f) * losses[-1]
    if loss < best_loss:
        best_loss = loss
    losses.append(loss)
    if loss > diverge_th * best_loss:
        print("Stopping early, the loss has diverged")
        break
    # reset model to initial parameters
    model.load_state_dict(torch.load('init_params.pt'))

    return lrs, losses

```

This is the result:



From this image, we can see that when $LR=10e-4$, loss decreases fastest. SO we choose $10e-4$ as the learning rate.

Training the model

This is the code of training:

```

def train(model, iterator, optimizer, criterion, device):
    epoch_loss = 0
    epoch_acc = 0
    model.train()
    for (x, y) in tqdm(iterator, desc="Training", leave=False):
        x = x.to(device)
        y = y.to(device)

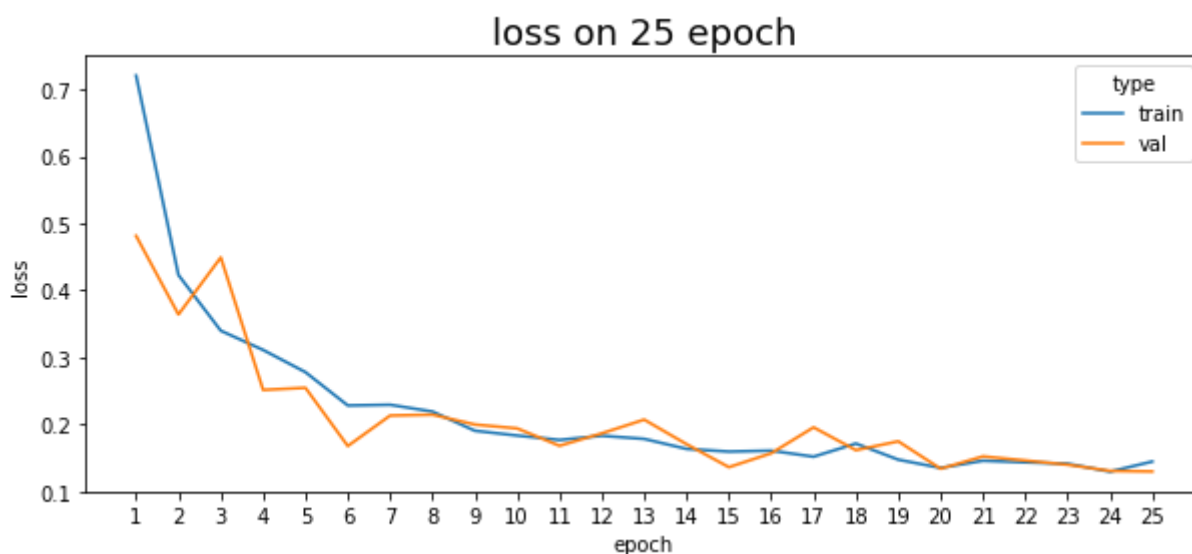
```

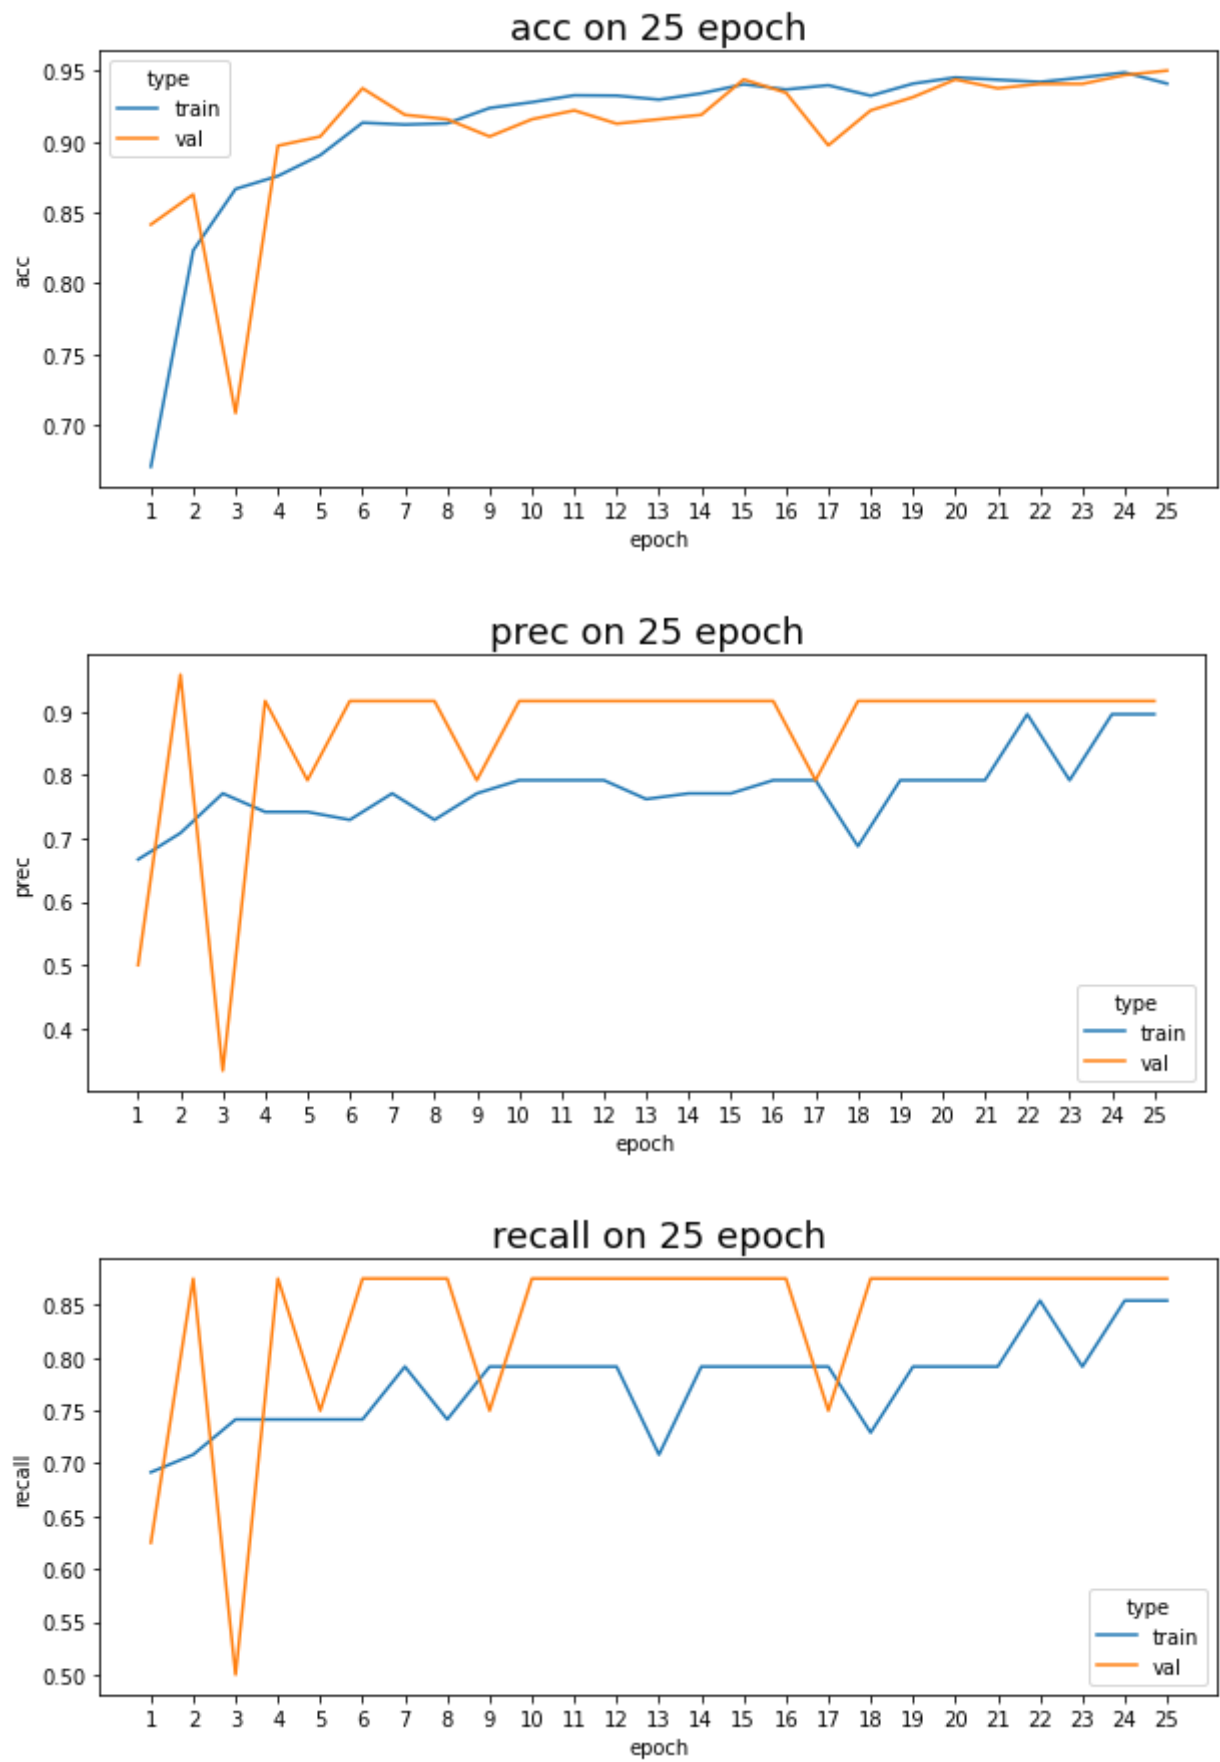
```
optimizer.zero_grad()
y_pred, _ = model(x)
loss = criterion(y_pred, y)
acc = calculate_accuracy(y_pred, y)
prec, recall = precision_recall(y_pred, y, average='macro',
num_classes=4)
loss.backward()
optimizer.step()
epoch_loss += loss.item()
epoch_acc += acc.item()
return epoch_loss / len(iterator), epoch_acc / len(iterator),
prec.to('cpu'), recall.to('cpu')
```

Model Performance Evaluation

After 25 epoches of training,

```
Epoch: 25 | Epoch Time: 0m 6s
Train Loss: 0.144 | Train Acc: 94.11% | Train Prec: 89.58% | Train
Recall: 85.42%
Val. Loss: 0.129 | Val. Acc: 95.03% | Val. Prec: 91.67% | Val.
Recall: 87.50%
```





And for test data, the prediction result is shown below.

