

The background image is a wide-angle landscape of a rugged mountain range. A prominent peak in the center-right is partially hidden by dark, billowing clouds. Below the mountain, a steep slope covered in sparse vegetation and rocky terrain descends towards the foreground. A thin, winding path or stream bed cuts through the slope. In the bottom right corner, there's a small cluster of buildings, possibly a hut or cabin, surrounded by a few people. The overall atmosphere is moody and dramatic.

# CISC3024 Pattern Recognition Project

by Zhang Huakang D-B9-2760-6



# Outline

-  **Data Loading** - loading data from file, build datasets and dataloaders
-  **Plot Example Image** - show some example image of each class
-  **Model Design** - *AlexNet*
-  **Model Training** - find the best lossing rate, and training the model
-  **Model Performance Evaluation** - built-in recording and camera view



# Data Loading

```
train_tran = transforms.Compose([
    transforms.RandomRotation(5),
    transforms.RandomHorizontalFlip(0.5),
    transforms.RandomCrop(64, padding=2),
    transforms.ToTensor(),
])
tran = transforms.Compose([
    transforms.ToTensor(),
])

class SatalliteDataset(Dataset):
    def __init__(self,images,labels, is_train):
        self.images_list=images
        self.labels_list=labels
        self.is_train=is_train

    def __len__(self):
        return len(self.labels_list)

    def __getitem__(self,idx):
        if self.is_train:
            return train_tran(self.images_list[idx]).to(device), self.labels_list[idx]
        return tran(self.images_list[idx]).to(device), self.labels_list[idx]
```

# Data Loading

```
images,labels = shuffle(loaddata(datapath,label_names), random_state=64)

X_train, X_Test, y_train, y_Test = train_test_split(images,labels, test_size=0.23, random_state=42)

X_val, X_test, y_val, y_test = train_test_split(X_Test, y_Test, test_size=0.23, random_state=42)

train_iterator = data.DataLoader(SatalliteDataset(X_train,y_train, True), batch_size=BATCH_SIZE)

valid_iterator = data.DataLoader(SatalliteDataset(X_test,y_test, False), batch_size=BATCH_SIZE)

test_iterator = data.DataLoader(SatalliteDataset(X_val,y_val, False), batch_size=BATCH_SIZE)
```

# Data Loading

```
images,labels = shuffle(loaddata(datapath,label_names), random_state=64)

X_train, X_Test, y_train, y_Test = train_test_split(images,labels, test_size=0.23, random_state=42)

X_val, X_test, y_val, y_test = train_test_split(X_Test, y_Test, test_size=0.23, random_state=42)

train_iterator = data.DataLoader(SatalliteDataset(X_train,y_train, True), batch_size=BATCH_SIZE)

valid_iterator = data.DataLoader(SatalliteDataset(X_test,y_test, False), batch_size=BATCH_SIZE)

test_iterator = data.DataLoader(SatalliteDataset(X_val,y_val, False), batch_size=BATCH_SIZE)
```

# Data Loading

```
images,labels = shuffle(loaddata(datapath,label_names), random_state=64)

X_train, X_Test, y_train, y_Test = train_test_split(images,labels, test_size=0.23, random_state=42)

X_val, X_test, y_val, y_test = train_test_split(X_Test, y_Test, test_size=0.23, random_state=42)

train_iterator = data.DataLoader(SatalliteDataset(X_train,y_train, True), batch_size=BATCH_SIZE)

valid_iterator = data.DataLoader(SatalliteDataset(X_test,y_test, False), batch_size=BATCH_SIZE)

test_iterator = data.DataLoader(SatalliteDataset(X_val,y_val, False), batch_size=BATCH_SIZE)
```

# Data Loading

```
images,labels = shuffle(loaddata(datapath,label_names), random_state=64)

X_train, X_Test, y_train, y_Test = train_test_split(images,labels, test_size=0.23, random_state=42)

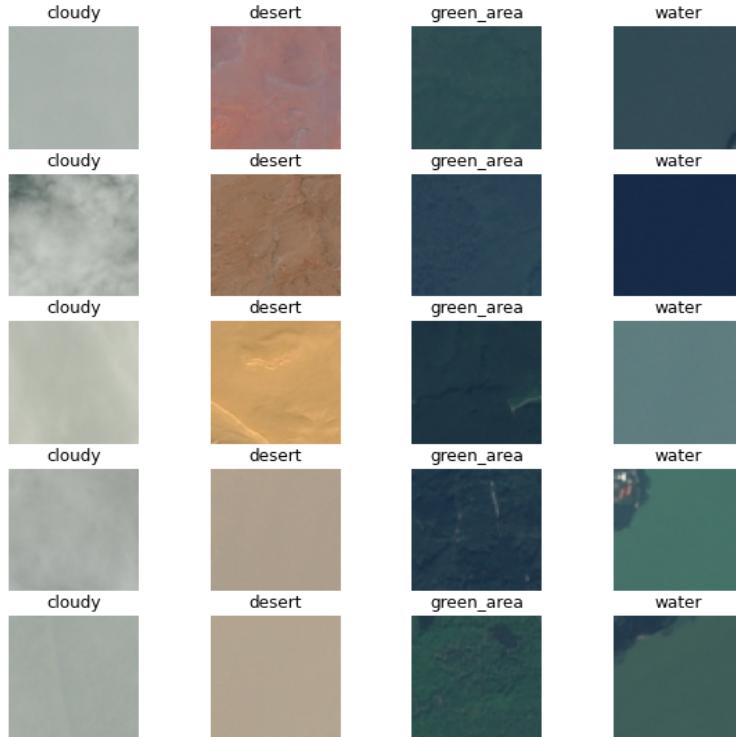
X_val, X_test, y_val, y_test = train_test_split(X_Test, y_Test, test_size=0.23, random_state=42)

train_iterator = data.DataLoader(SatalliteDataset(X_train,y_train, True), batch_size=BATCH_SIZE)

valid_iterator = data.DataLoader(SatalliteDataset(X_test,y_test, False), batch_size=BATCH_SIZE)

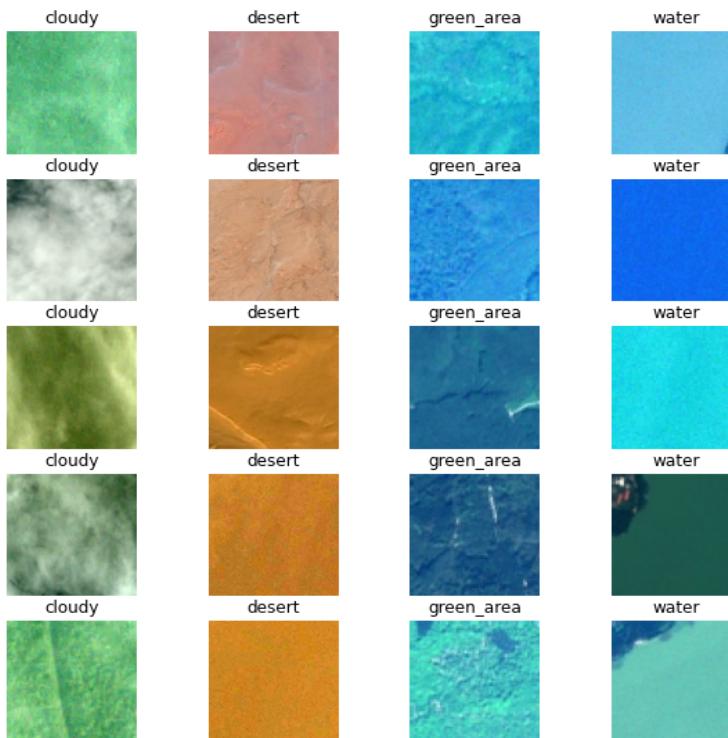
test_iterator = data.DataLoader(SatalliteDataset(X_val,y_val, False), batch_size=BATCH_SIZE)
```

# Dataset Example



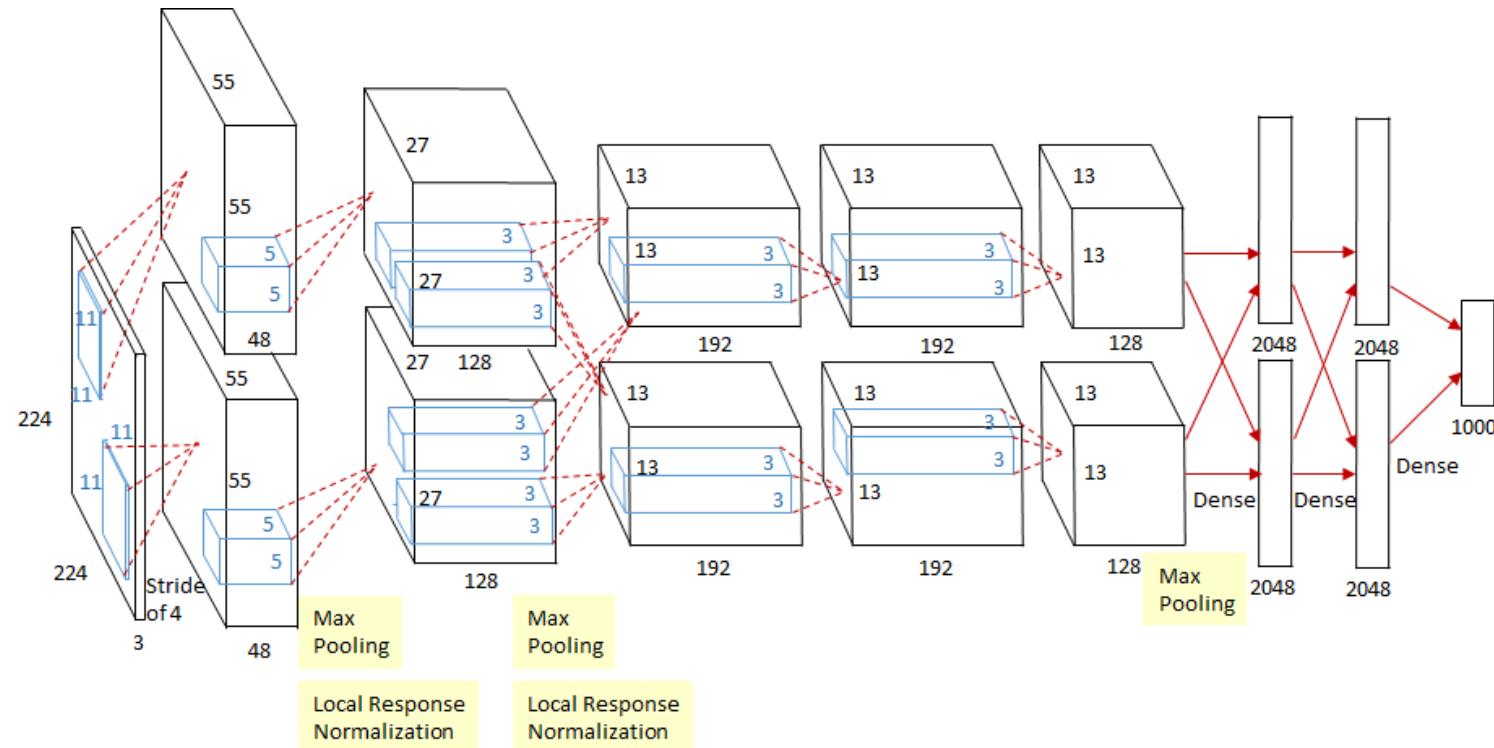
# Dataset Example

with *normalization*



# Model Design

AlexNet



# Model Design

features

```
self.features = nn.Sequential(  
    nn.Conv2d(3, 64, 3, 2, 1),  
    nn.MaxPool2d(2),  
    nn.ReLU(inplace=True),  
    nn.Conv2d(64, 192, 3, padding=1),  
    nn.MaxPool2d(2),  
    nn.ReLU(inplace=True),  
    nn.Conv2d(192, 384, 3, padding=1),  
    nn.ReLU(inplace=True),  
    nn.Conv2d(384, 256, 3, padding=1),  
    nn.ReLU(inplace=True),  
    nn.Conv2d(256, 256, 3, padding=1),  
    nn.MaxPool2d(2),  
    nn.ReLU(inplace=True)  
)
```

# Model Design

classifier

```
self.classifier = nn.Sequential(  
    nn.Dropout(0.5),  
    nn.Linear(4096, 4096),  
    nn.ReLU(inplace=True),  
    nn.Dropout(0.5),  
    nn.Linear(4096, 4096),  
    nn.ReLU(inplace=True),  
    nn.Linear(4096, output_dim),  
)
```

The model has **35,830,596** trainable parameters

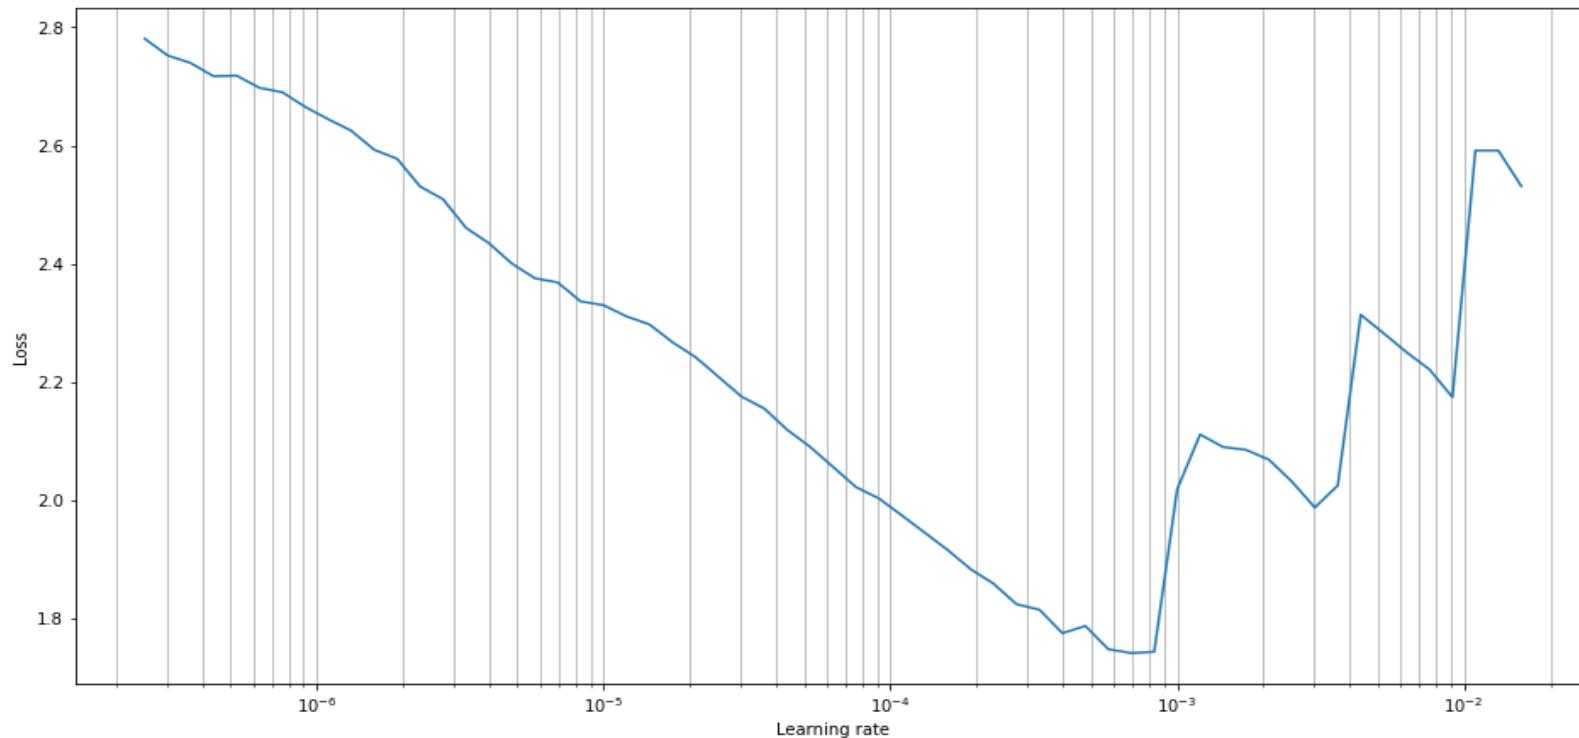
# Model Training

Get the best Learning Rate, in `LRFinder` class

```
def range_test(self, iterator, end_lr=10, num_iter=100, smooth_f=0.05, diverge_th=5):
    lrs = []
    losses = []
    best_loss = float('inf')
    lr_scheduler = ExponentialLR(self.optimizer, end_lr, num_iter)
    iterator = IteratorWrapper(iterator)
    for iteration in range(num_iter):
        loss = self._train_batch(iterator)
        lrs.append(lr_scheduler.get_last_lr()[0])
        lr_scheduler.step()
        if iteration > 0:
            loss = smooth_f * loss + (1 - smooth_f) * losses[-1]
        if loss < best_loss:
            best_loss = loss
        losses.append(loss)
        if loss > diverge_th * best_loss:
            print("Stopping early, the loss has diverged")
            break
    model.load_state_dict(torch.load('init_params.pt'))
    return lrs, losses
```

# Model Training

FOUND\_LR = 1e-4



# Model Training

```
for epoch in trange(EPOCHS, desc="Epochs"):
    start_time = time.monotonic()
    train_loss, train_acc = train(model, train_iterator, optimizer, criterion, device)
    valid_loss, valid_acc = evaluate(model, valid_iterator, criterion, device)
    end_time = time.monotonic()
    epoch_mins, epoch_secs = epoch_time(start_time, end_time)
    print(f'Epoch: {epoch+1:02} | Epoch Time: {epoch_mins}m {epoch_secs}s')
    print(f'\tTrain Loss: {train_loss:.3f} | Train Acc: {train_acc*100:.2f}%')
    print(f'\tVal. Loss: {valid_loss:.3f} | Val. Acc: {valid_acc*100:.2f}%')
```

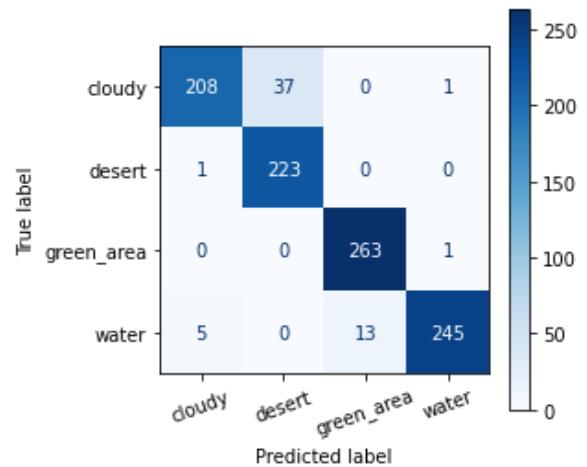
# Model Performance Evaluation

Epoch: 25 | Epoch Time: 0m 5s

Train Loss: 0.116 | Train Acc: 95.74%

Val. Loss: 0.193 | Val. Acc: 92.22%

Test Loss: 0.161 | Test Acc: 93.81%



# Q&A